

movl Src, Dest	Dest = Src	
movsbl Src, Dest	Dest (long) = Src (byte), sign extend	
addl Src, Dest	Dest = Dest + Src	
subl Src, Dest	Dest = Dest - Src	
imull Src, Dest	Dest = Dest * Src	
sall Src, Dest	Dest = Dest << Src	
sarl Src, Dest	Dest = Dest >> Src	arithmetic shift
shrl Src, Dest	Dest = Dest >> Src	logical shift
xorl Src, Dest	Dest = Dest ^ Src	
andl Src, Dest	Dest = Dest & Src	
orl Src, Dest	Dest = Dest   Src	
incl Dest	Dest = Dest + 1	
decl Dest	Dest = Dest - 1	
negl Dest	Dest = - Dest	
notl Dest	Dest = ~ Dest	
leal Src, Dest	Dest = address of Src	
cmpl Src2, Src1	Sets CCs Src1 - Src2	
testl Src2, Src1	Sets CCs Src1 & Src2	

Hex digit	0	1	2	3	4	5	6	7
Decimal value	0	1	2	3	4	5	6	7
Binary value	0000	0001	0010	0011	0100	0101	0110	0111
Hex digit	8	9	A	B	C	D	E	F
Decimal value	8	9	10	11	12	13	14	15
Binary value	1000	1001	1010	1011	1100	1101	1110	1111

Figure 2.2 Hexadecimal notation. Each Hex digit encodes one of 16 values.

	~	&	0	1		0	1	^	0	1
0	1	0	0	0	0	0	1	0	0	1
1	0	1	0	1	1	1	1	1	1	0

Figure 2.7 Operations of Boolean algebra. Binary values 1 and 0 encode logic values TRUE and FALSE, while operations ~, &, |, and ^ encode logical operations NOT, AND, OR, and EXCLUSIVE-OR, respectively.

### IEEE Floating Point Representation

	sign	exp	mantissa	bias	
single	1	8	23	127	(32-bit)
double	1	11	52	1023	(64-bit)

bias =  $2^{(k-1)} - 1$  where k = number of bits in exp

- Normalized Exponent (E) = exp - bias

S	≠ 0 & ≠ 255	f
---	-------------	---

- Denormalized Exponent (E) = 1 - bias

S	0 0 0 0 0 0 0 0	f
---	-----------------	---

3a. Infinity

S	1 1 1 1 1 1 1 1	0 0
---	-----------------	---

3b. NaN

S	1 1 1 1 1 1 1 1	≠ 0
---	-----------------	-----

Figure 2.32 Categories of single-precision, floating-point values. The value of the exponent determines whether the number is (1) normalized, (2) denormalized, or a (3) special value.

<b>Example</b>	sign	exp	mantissa
-13.9 (base 10) in 8 bits	1	0 0 1 1	1 0 1
-13 (base 10) = ? (base 2)		0.9 (base 10) = ? (base 2)	
13/2 = 6 r1		0.9 * 2 =	1.8 1 .8
6/2 = 3 r0		0.8 * 2 =	1.6 1 .6
3/2 = 1 r1		0.6 * 2 =	1.2 1 .2
1/2 = 0 r1		0.2 * 2 =	0.4 0 .4
		0.4 * 2 =	0.8 0 .8
		(Stop when repeat begins)	
1101		111100	
-13.9 in binary =	-1101.11100	* 2^-0	
	1.10111100	* 2^-3	
	mantissa	exp	

Instruction	Synonym	Effect	Set condition
sete D	setz	D ← ZF	Equal / zero
setne D	setnz	D ← ~ZF	Not equal / not zero
sets D		D ← SF	Negative
setns D		D ← ~SF	Nonnegative
setg D	setnle	D ← ~(SF ^ OF) & ~ZF	Greater (signed >)
setge D	setnl	D ← ~(SF ^ OF)	Greater or equal (signed >=)
setl D	setnge	D ← SF ^ OF	Less (signed <)
setle D	setng	D ← (SF ^ OF)   ZF	Less or equal (signed <=)
seta D	setnbe	D ← ~CF & ~ZF	Above (unsigned >)
setae D	setnb	D ← ~CF	Above or equal (unsigned >=)
setb D	setnae	D ← CF	Below (unsigned <)
setbe D	setna	D ← CF   ZF	Below or equal (unsigned <=)

Figure 3.11 The SET instructions. Each instruction sets a single byte to 0 or 1 based on some combination of the condition codes. Some instructions have "synonyms," i.e., alternate names for the same machine instruction.

Instruction	Based on	Description
CMP	S <sub>2</sub> , S <sub>1</sub>	S <sub>1</sub> - S <sub>2</sub> Compare
cmnb		Compare byte
cmpl		Compare word
cmpl		Compare double word
TEST	S <sub>2</sub> , S <sub>1</sub>	S <sub>1</sub> & S <sub>2</sub> Test
testb		Test byte
testw		Test word
testl		Test double word

Figure 3.10 Comparison and test instructions. These instructions set the condition codes without updating any other registers.

movl Val, Val	jne label	jump not equal
val: constant integer value	js label	jump negative
movl \$17, %eax	jns label	jump non-negative
	jg label	jump greater (signed)
	jge label	jump greater or equal (signed)
	jl label	jump less (signed)
	jle label	jump less or equal (signed)
	ja label	jump above (unsigned)
	jnb label	jump below (unsigned)
	push Src	%esp = %esp - 4, Mem[%esp] = Src
	pop Dest	Dest = Mem[%esp], %esp = %esp + 4
	call label	push address of next instruction, jmp label
	ret	%ip = Mem[%esp], %esp = %esp + 4

<b>Endianness</b>	<b>Bit Shifting</b>
Little Endian: Least significant digit comes first	Operation
Big endian: Least significant digit comes last	Argument x
For example: 0xFFAE07	x << 4
Little endian: 07 AE FF	x >> 4 (logical)
Big endian: FF AE 07	x >> 4 (arithmetic)
void inplace_swap(int *x, int *y) {	
2 *y = *x ^ *y; /* Step 1 */	
3 *x = *x ^ *y; /* Step 2 */	
4 *y = *x ^ *y; /* Step 3 */	
5 }	

<b>Multiplication by Bitwise Shifting Left</b>	<b>Division by Powers of Two Using Bitwise Shifting Right</b>
x*14	x/2^k
14 = 2^3 + 2^2 + 2^1	(x << 0 ? x + (x << k) - 1 : x) >> k
(x << 3) + (x << 2) + (x << 1)	
14 = 2^4 - 2^1	
(x << 4) - (x << 1)	


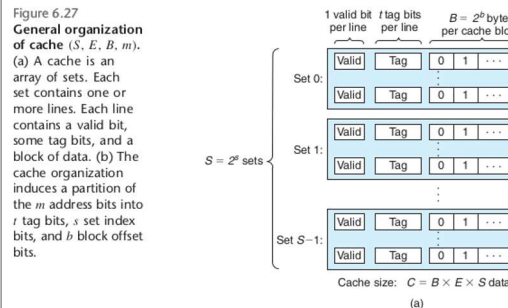


Figure 3.24 Stack frames for caller and swap\_add. Procedure swap\_add retrieves its arguments from the stack frame for caller.



b	byte
w	word (2 bytes)
l	long (4 bytes)

<b>Condition codes</b>
CF Carry Flag
ZF Zero Flag
SF Sign Flag
OF Overflow Flag

<b>Registers</b>
%eax
%ecx
%edx
%ebx
%esi
%edi
%esp
%ebp

<b>Amdahl's Law</b>
Increase in speed (performance) = $\frac{1}{\text{fraction enhanced} \cdot \text{speedup enhancement} + (1 - \text{fraction enhanced})}$
where 'fraction enhanced' is the fractional part of the thing enhanced
<b>Examples</b>
Speedup of 80x with 100 processor cores
$80 = \frac{1}{(\text{frac parallel} / 100) + (1 - \text{frac parallel})}$
$\Rightarrow 0.8 * \text{frac parallel} + 80(1 - \text{frac parallel}) = 1$
frac parallel = $(80 - 1) / 79.2 = 0.9975$ 99% of the project will have to change to benefit from parallelism
A trucker drives 2500 km averaging 100 km/hr (25 hours)
If for 1500 km of the trip the trucker traveled 150 km/hr, what will be the trucker's speedup for the trip?
enhancement = $\frac{1}{(150 / 100) + (1 - 1.5) / 0.6} = 1.25$ times faster
fraction enhanced = $(1500 \text{ km} / 2500 \text{ km}) = 0.6$
How fast would the trucker have to drive over the portion of 1500 km of the trip to get an overall speedup of 5/3?
$\frac{5}{3} = \frac{1}{(0.6 / \text{enhancement}) + (1 - 0.6)}$ enhancement = 3.0 times as fast $\Rightarrow$ 300 km/hr

<b>Figure 6.27</b>	<b>Figure 6.28</b>
<b>General organization of cache</b> (S, E, B, m).	<b>Summary of cache parameters.</b>
(a) A cache is an array of sets. Each set contains one or more lines. Each line contains a valid bit, some tag bits, and a block of data.	<b>Fundamental parameters</b>
(b) The cache organization induces a partitioning of the m address bits into t tag bits, s set index bits, and b block offset bits.	<b>Derived quantities</b>
	Parameter Description
	M = 2^m Maximum number of unique memory addresses
	S = log <sub>2</sub> (S) Number of set index bits
	E Number of lines per set
	B = 2^b Block size (bytes)
	m = log <sub>2</sub> (M) Number of physical (main memory) address bits
	Parameter Description
	M = 2^m Maximum number of unique memory addresses
	S = log <sub>2</sub> (S) Number of set index bits
	B = log <sub>2</sub> (B) Number of block offset bits
	t = m - (s + b) Number of tag bits
	C = B * E * S Cache size (bytes) not including overhead such as the valid and tag bits

Figure 6.28 Summary of cache parameters.